



# Machine Learning

## LABORATORY: CNN Homework

NAME:

STUDENT ID#:

### Objectives:

- In this assignment, you will train a Convolutional Neural Network (CNN) to classify images of cats and dogs using a dataset and starter template provided to you.
- The template code includes a simple CNN model that serves as a baseline. However, this baseline model does not achieve high accuracy.
- Your goal is to improve the model's performance by modifying and fine-tuning the CNN architecture.
- You are encouraged to: Adjust the layer, add data augmentation and apply regularization techniques such as Dropout, BatchNormalization, or custom pooling strategies.

### Part 1. Instruction

- The template code includes a simple CNN model that serves as a baseline. However, this baseline model does not achieve high accuracy.
- Your task is to improve the model's performance by modifying and fine-tuning the CNN architecture.
- You are encouraged to:
  - Adjust layer configurations (number of layers, filters, kernel sizes, activations, etc.)
  - Add data augmentation to improve generalization
  - Apply regularization techniques such as Dropout, BatchNormalization, or custom pooling strategies
- **However**, you are not allowed to use pre-trained models such as VGG, ResNet, EfficientNet, etc. (No transfer learning or model imports from keras.applications).
- Your model should be designed, trained, and fine-tuned by yourself, from scratch using Keras/TensorFlow layers.
- If you're not familiar with TensorFlow or Keras, check these links:
  - Tensorflow Setup Guide <https://github.com/JTKostman/keras-tensorflow-windows-installation>
  - Keras Quick Start : [https://keras.io/getting\\_started/](https://keras.io/getting_started/)
- If you're using your **local PC**, you may need to install TensorFlow and Keras manually. If you're using **Google Colab**, they're already installed — but be aware that Colab has **limited training time**.
- Dataset: <https://drive.google.com/file/d/1eyrNGFIM83pf-TETo30Cvjm7kYuCFAqu/view?usp=sharing> (Also included in the template code for easy loading)



## Part 2. Code Template

| Step | Procedure   |
|------|---|
| 1    | <pre> import pandas as pd import IPython.display as display import tensorflow as tf from tensorflow.keras import layers import numpy as np import os, random import matplotlib.pyplot as plt print(tf.__version__) </pre>   |
| 2    | <pre> # ===== Download Dataset ===== #!pip install -U gdown  import gdown, zipfile  file_id = "1eyrNGF1M83pf-TETo30Cvjm7kYuCFAqu" # &lt;-- from your real zip file link gdown.download(f"https://drive.google.com/uc?id={file_id}", output="dataset.zip", quiet=False)  with zipfile.ZipFile("dataset.zip", 'r') as zip_ref:     zip_ref.extractall("dataset") </pre>   |
| 3    | <pre> # =====Dataset Preparation ===== IMAGE_HEIGHT=128 IMAGE_WIDTH=128 BATCH_SIZE=64 def get_pathframe(path):     '''     Get all the images paths and its corresponding labels     Store them in pandas dataframe     '''     filenames = os.listdir(path)     categories = []     paths=[]     for filename in filenames:         paths.append(path+filename)         category = filename.split('.')[0]         if category == 'dog':             categories.append(1)         else:             categories.append(0)  df= pd.DataFrame({ </pre> |



|   |  |
|---|--|
|   | <pre>         'filename': filenames,         'category': categories,         'paths':paths     })     return df  df=get_pathframe("dataset/dataset/dataset/") df.tail(5) </pre>  |
| 4 | <pre> # ===== Convert to tensor ===== def load_and_preprocess_image(path):     '''     Load each image and resize it to desired shape     '''     image = tf.io.read_file(path)     image = tf.image.decode_jpeg(image, channels=3)     image = tf.image.resize(image, [IMAGE_WIDTH, IMAGE_HEIGHT])     image /= 255.0 # normalize to [0,1] range     return image  def convert_to_tensor(df):     '''     Convert each data and labels to tensor     '''     path_ds = tf.data.Dataset.from_tensor_slices(df['paths'])     image_ds = path_ds.map(load_and_preprocess_image)     # onehot_label=tf.one_hot(tf.cast(df['category'], tf.int64),2) if     using softmax     onehot_label=tf.cast(df['category'], tf.int64)     label_ds = tf.data.Dataset.from_tensor_slices(onehot_label)      return image_ds,label_ds  X,Y=convert_to_tensor(df) print("Shape of X in data:", X) print("Shape of Y in data:", Y) </pre> |
| 5 | <pre> #Plot Images dataset=tf.data.Dataset.zip((X,Y)).shuffle(buffer_size=2000) dataset_train=dataset.take(22500) dataset_test=dataset.skip(22500) </pre>  |



|   |  |
|---|--|
|   | <pre> dataset_train=dataset_train.batch(BATCH_SIZE, drop_remainder=True) dataset_test=dataset_test.batch(BATCH_SIZE, drop_remainder=True) dataset_train  def plotimages(imagesls):     fig, axes = plt.subplots(1, 5, figsize=(20,20))     axes = axes.flatten()     for image,ax in zip(imagesls, axes):         ax.imshow(image)         ax.axis('off')  imagesls=[] for n, image in enumerate(X.take(5)):     imagesls.append(image)  plotimages(imagesls) </pre>   |
| 6 | <pre> #Model Design def My_CNNmodel():      model = tf.keras.models.Sequential()     model.add(layers.Conv2D(8, (3, 3), padding='same',activation='relu', input_shape=(IMAGE_WIDTH, IMAGE_HEIGHT, 3)))     model.add(layers.MaxPooling2D(pool_size=(2,2)))     model.add(layers.Conv2D(16, (3, 3), padding='same',activation='relu'))     model.add(layers.MaxPooling2D(pool_size=(2,2)))     model.add(layers.Conv2D(32, (3, 3), padding='same',activation='relu'))     model.add(layers.MaxPooling2D(pool_size=(2,2)))     model.add(layers.Conv2D(64, (3, 3), padding='same',activation='relu'))     model.add(layers.MaxPooling2D(pool_size=(2,2)))     model.add(layers.Flatten())     model.add(layers.Dense(512, activation='relu'))     model.add(layers.Dense(1, activation='sigmoid'))      opt=tf.keras.optimizers.Adam(0.001)     model.compile(optimizer=opt,                     loss='binary_crossentropy', # loss='categorical_crossentropy' if softmax                     metrics=['accuracy']) </pre> |



|   |  |
|---|--|
|   | <pre> return model  model=My_CNNmodel() model.summary() </pre>   |
| 7 | <pre> #Training the model #You can adjust the epochs to get better training results, be aware with overfitting  hist=model.fit(dataset_train,epochs=2,validation_data=dataset_test) #Save trained model model.save("student_ID.keras") #Save the model with your student ID </pre>   |
|   | <pre> #Plot training results def plot_model_history(model_history, acc='accuracy', val_acc='val_accuracy'):     fig, axs = plt.subplots(1, 2, figsize=(15, 5))      # Accuracy plot     axs[0].plot(range(1, len(model_history.history[acc]) + 1), model_history.history[acc])     axs[0].plot(range(1, len(model_history.history[val_acc]) + 1), model_history.history[val_acc])     axs[0].set_title('Model Accuracy')     axs[0].set_ylabel('Accuracy')     axs[0].set_xlabel('Epoch')     axs[0].set_xticks(np.arange(1, len(model_history.history[acc]) + 1))     axs[0].legend(['train', 'val'], loc='best')      # Loss plot     axs[1].plot(range(1, len(model_history.history['loss']) + 1), model_history.history['loss'])     axs[1].plot(range(1, len(model_history.history['val_loss']) + 1), model_history.history['val_loss'])     axs[1].set_title('Model Loss')     axs[1].set_ylabel('Loss')     axs[1].set_xlabel('Epoch')     axs[1].set_xticks(np.arange(1, len(model_history.history['loss']) + 1))     axs[1].legend(['train', 'val'], loc='best') </pre> |

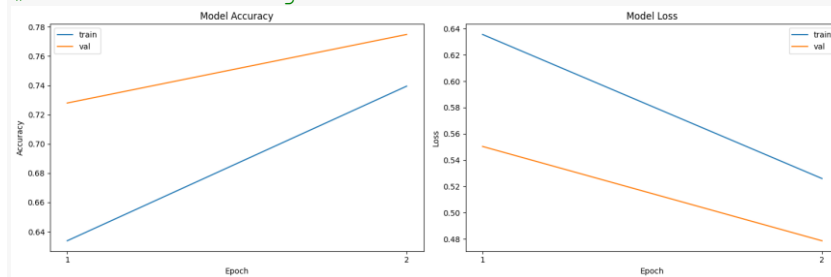


|  |  |
|--|--|
|  | <pre> plt.tight_layout() plt.show() plot_model_history(hist) </pre>  |
|  | <pre> #Evaluate the model import seaborn as sns from sklearn.metrics import confusion_matrix, classification_report import matplotlib.pyplot as plt import numpy as np  # Evaluate the model loss, accuracy = model.evaluate(dataset_test) print("Test: accuracy = %f ; loss = %f " % (accuracy, loss))  # Predict values y_pred = model.predict(dataset_test) y_p = np.where(y_pred &gt; 0.5, 1, 0) # for binary classification  # Extract ground truth labels test_data = dataset_test.unbatch() y_g = [] for image, label in test_data:     y_g.append(label.numpy())  # Convert to flat array if needed y_g = np.array(y_g).flatten() y_p = y_p.flatten()  # Compute confusion matrix confusion_mtx = confusion_matrix(y_g, y_p)  # Plot f, ax = plt.subplots(figsize=(8, 8)) sns.heatmap(confusion_mtx, annot=True, linewidths=0.01, cmap="Blues", linecolor="gray", fmt='.1f', ax=ax) plt.xlabel("Predicted Label") plt.ylabel("True Label") plt.title("Confusion Matrix") plt.show() </pre> |
|  | <pre> # Generate a classification report report = classification_report(y_g, y_p, target_names=['0','1']) print(report) </pre>   |



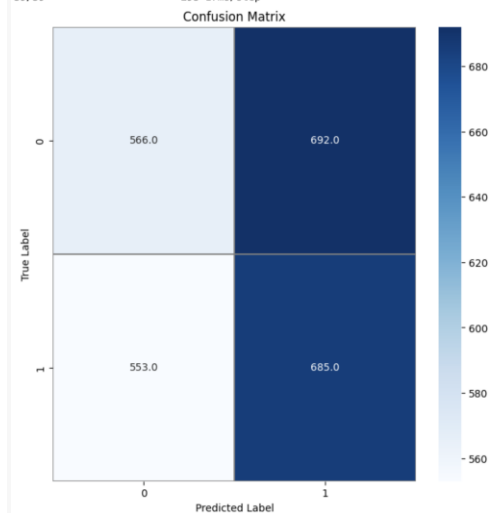
## #Example Output (For references only)

### #1. Plot Training Results



### #2. Confusion Matrix

39/39 ————— 19s 22ms/step - accuracy: 0.7777 - loss: 0.4819  
 Test: accuracy = 0.770433 ; loss = 0.482054  
 39/39 ————— 19s 17ms/step



### #3. Classification Report

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.51      | 0.45   | 0.48     | 1258    |
| 1            | 0.50      | 0.55   | 0.52     | 1238    |
| accuracy     |           |        | 0.50     | 2496    |
| macro avg    | 0.50      | 0.50   | 0.50     | 2496    |
| weighted avg | 0.50      | 0.50   | 0.50     | 2496    |

### #4. Model Summary

Model: "sequential"

| Layer (type)                   | Output Shape        | Param #   |
|--------------------------------|---------------------|-----------|
| conv2d (Conv2D)                | (None, 128, 128, 8) | 224       |
| max_pooling2d (MaxPooling2D)   | (None, 64, 64, 8)   | 0         |
| conv2d_1 (Conv2D)              | (None, 64, 64, 16)  | 1,168     |
| max_pooling2d_1 (MaxPooling2D) | (None, 32, 32, 16)  | 0         |
| conv2d_2 (Conv2D)              | (None, 32, 32, 32)  | 4,640     |
| max_pooling2d_2 (MaxPooling2D) | (None, 16, 16, 32)  | 0         |
| conv2d_3 (Conv2D)              | (None, 16, 16, 64)  | 18,496    |
| max_pooling2d_3 (MaxPooling2D) | (None, 8, 8, 64)    | 0         |
| flatten (Flatten)              | (None, 4096)        | 0         |
| dense (Dense)                  | (None, 512)         | 2,097,664 |
| dense_1 (Dense)                | (None, 1)           | 513       |

Total params: 2,122,705 (8.10 MB)  
 Trainable params: 2,122,705 (8.10 MB)  
 Non-trainable params: 0 (0.00 B)



## Grading Assignment & Submission (70%)

### Implementation(45%):

1. (5%) **Fine-Tuning:** Adjust model parameters to achieve better accuracy
2. (20%) **Model redesign with techniques:** Redesign the architecture using advanced techniques such as data augmentation, Dropout, BatchNormalization, or self-designed structures (should achieve better accuracy), brief **explain your model summary in the report.**
3. (15%) **Accuracy Improvement:** Demonstrate significantly improved validation accuracy compared to the baseline model. **Explain your performance improvement.**
4. (5%) **Include all evaluation results:** model summary, training plots, confusion matrix, and classification report.

### Question(25%):

5. (7%) What changes did you make to the CNN architecture? Why did you choose those modifications, and how did they affect model performance?
6. (8%) What methods did you use to reduce overfitting (if any)? How effective were they? Explain with reference to training/validation curves.
7. (10%) Based on your confusion matrix, what kinds of misclassifications occurred most frequently? What might be causing these errors? How would you attempt to reduce these errors?

### Submission :

1. Report: Provide your screenshots of your results in the last pages of this PDF File.
2. Code: Submit your complete Python script in either .py or .ipynb format.
3. Upload both your report and code to the E3 system (**Labs6 Homework Assignment**). Name your files correctly:
  - a. Report: StudentID\_Lab6\_Homework.pdf
  - b. Code: StudentID\_Lab6\_Homework.py or StudentID\_Lab6\_Homework.ipynb
4. Deadline: 16:20 PM
5. Plagiarism is **strictly prohibited**. Submitting copied work from other students will result in penalties.

### Results and Discussion:

